

Learning From Experiments With Causal Machine Learning

A case study using **metalearners**

Francesc Martí Escofet (@fmartiescofet)


Kevin Klein (@kevkle)



We live in a budget-constrained world.


Which students should be coached?

National Study of Learning Mindsets




Before data
collection

National Study of Learning Mindsets




Before data
collection



Ninth grade

National Study of Learning Mindsets



The data, more formally

Every datapoint i corresponds to a student.

Name	Symbol	Definition
Covariates	X_i	Properties of the student or the student's school
Treatment assignments	W_i	$\begin{cases} 1 & \text{if received coaching} \\ 0 & \text{if didn't receive coaching} \end{cases}$
Outcome	Y_i	GPA ($\in \mathbb{R}$) after treatment

$$\mathcal{D} = \{(X_i, W_i, Y_i)\}$$

The data, the details

- $n = 10,391$
 - $\sim 1/3$ received coaching
- Originally from National Study of Learning Mindsets
 - Nature, September 2019
- We used an anonymized version from [Athey and Wager](#)
 - All continuous features have been transformed to a standard Normal

The data, an excerpt

	schoolid	success_expect	ethnicity	gender	frst_in_family	school_urban
3625	75	5	2	1	1	
3037	5	2	2	1	1	
2574	17	7	4	2	1	
1488	47	6	5	2	1	
3677	74	6	11	2	1	

Loading the data


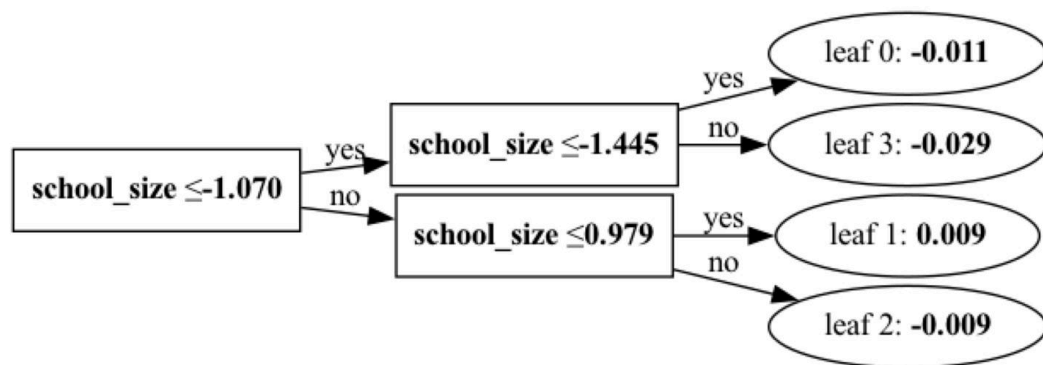
```
import pandas as pd

df = pd.read_csv("learning_mindset.csv")
```

```
categorical_feature_columns = [
    "ethnicity",
    "gender",
    "first_in_family",
    "school_urbanicity",
    "schoolid",
]
```

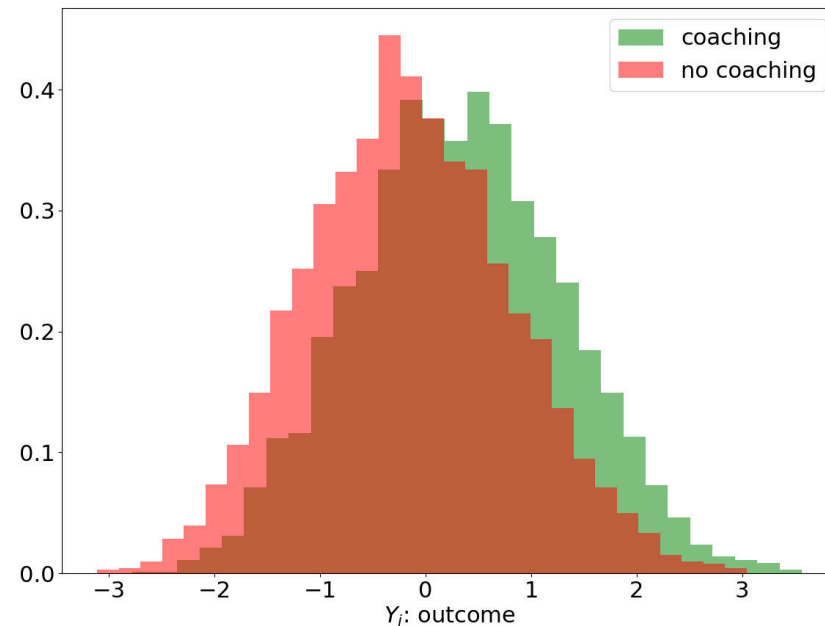

```
for column in categorical_feature_columns:
    df[column] = df[column].astype("category")
```

Why categoricals?



```
ax.hist(df[treatment_column])
```

```
ax.hist(df[W=1][outcome_column], density=True)  
ax.hist(df[W=0][outcome_column], density=True)
```



What do we do with the data now?

- Remember, our original question was
 - (A) Which students should be coached?
- We'll reduce said question to the following question
 - (B) How much would a student like student i profit from a growth mindset coaching?

What do we do with the data now?

- (B) How much would a student like student i profit from a growth mindset coaching?

formalism: **Conditional Average Treatment Effect (CATE)**

$$\tau(X_i) = \mathbb{E}[Y(\text{coaching}) - Y(\text{no coaching}) | X = X_i]$$

What do we do with the data now?

- (B) How much would a student like student i profit from a growth mindset coaching?

formalism: **Conditional Average Treatment Effect** (CATE)

$$\tau(X_i) = \mathbb{E}[\underbrace{Y(\text{coaching}) - Y(\text{no coaching})}_{\text{treatment effect}} | X = X_i]$$

What do we do with the data now?

- (B) How much would a student like student i profit from a growth mindset coaching?

formalism: **Conditional Average Treatment Effect (CATE)**

$$\tau(X_i) = \underbrace{\mathbb{E}}_{\text{average}} [Y(\text{coaching}) - Y(\text{no coaching}) | X = X_i]$$

What do we do with the data now?

- (B) How much would a student like student i profit from a growth mindset coaching?

formalism: **Conditional Average Treatment Effect (CATE)**

$$\tau(X_i) = \mathbb{E}[Y(\text{coaching}) - Y(\text{no coaching}) | \underbrace{X = X_i}_{\text{conditional}}]$$

What do we do with the data now?

- (B) How much would a student like student i profit from a growth mindset coaching?

formalism: **Conditional Average Treatment Effect (CATE)**

$$\tau(X_i) = \mathbb{E}[Y(\text{coaching}) - Y(\text{no coaching}) | X = X_i]$$


- (A) Which students should be coached?

formalism: **policy**

$$\pi(X_i) = \begin{cases} 1 & \text{if } \hat{\tau}(X_i) \geq c_{budget} \\ 0 & \text{otherwise} \end{cases}$$


MetaLearners

- MetaLearners are **CATE models** which rely on typical, **arbitrary machine learning estimators** (classifiers or regressors) as **components**.
- Some examples include the S-Learner, T-Learner, F-Learner, X-Learner, R-Learner, M-Learner and DR-Learner.



MetaLearners

- Input
 - W : Treatment assignments
 - X : Covariates/features
 - Y : Outcomes
- Output
 - $\hat{\tau}(X)$: CATE estimates



Creating a first MetaLearner

```
from metalearners import RLearner
from lightgbm import LGBMRegressor, LGBMClassifier
```

```
rlearner = RLearner(
    nuisance_model_factory=LGBMRegressor,
    propensity_model_factory=LGBMClassifier,
    treatment_model_factory=LGBMRegressor,
    is_classification=False,
    n_variants=2,
)
```

Creating a first MetaLearner

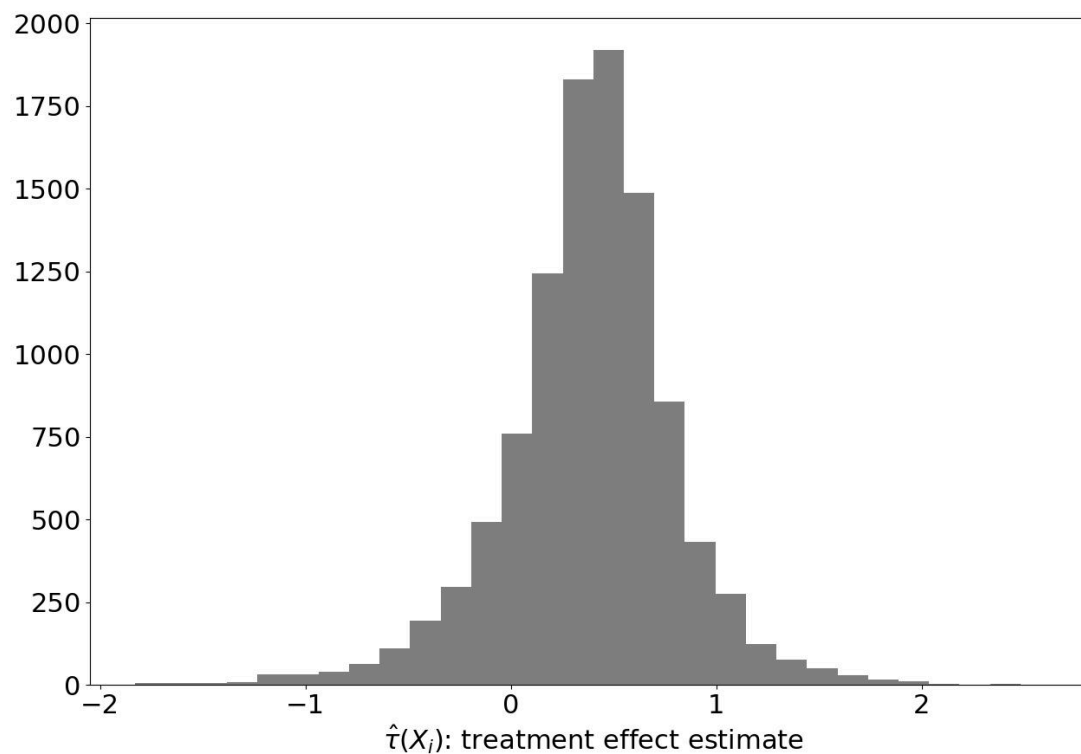
```
from metalearners import RLearner
from lightgbm import LGBMRegressor, LGBMClassifier
```

```
rlearner = RLearner(
    nuisance_model_factory=LGBMRegressor,
    propensity_model_factory=LGBMClassifier,
    treatment_model_factory=LGBMRegressor,
    is_classification=False,
    n_variants=2,
)
```

```
rlearner.fit(
    X=df[feature_columns], y=df[outcome_column], w=df[treatment_column]
)
```

Predicting with a MetaLearner

```
rlearner.predict(df[feature_columns], is_oos=False)
```






Hyperparameter optimization


- HPO can have massive impacts on the prediction quality in regular Machine Learning
- According to [Machlanski et. al \(2023\)](#) this also happens in MetaLearners
- Three levels to optimize for:
 - The MetaLearner architecture
 - The model to choose per base estimator
 - The model hyperparameters per base model




Performing a grid search



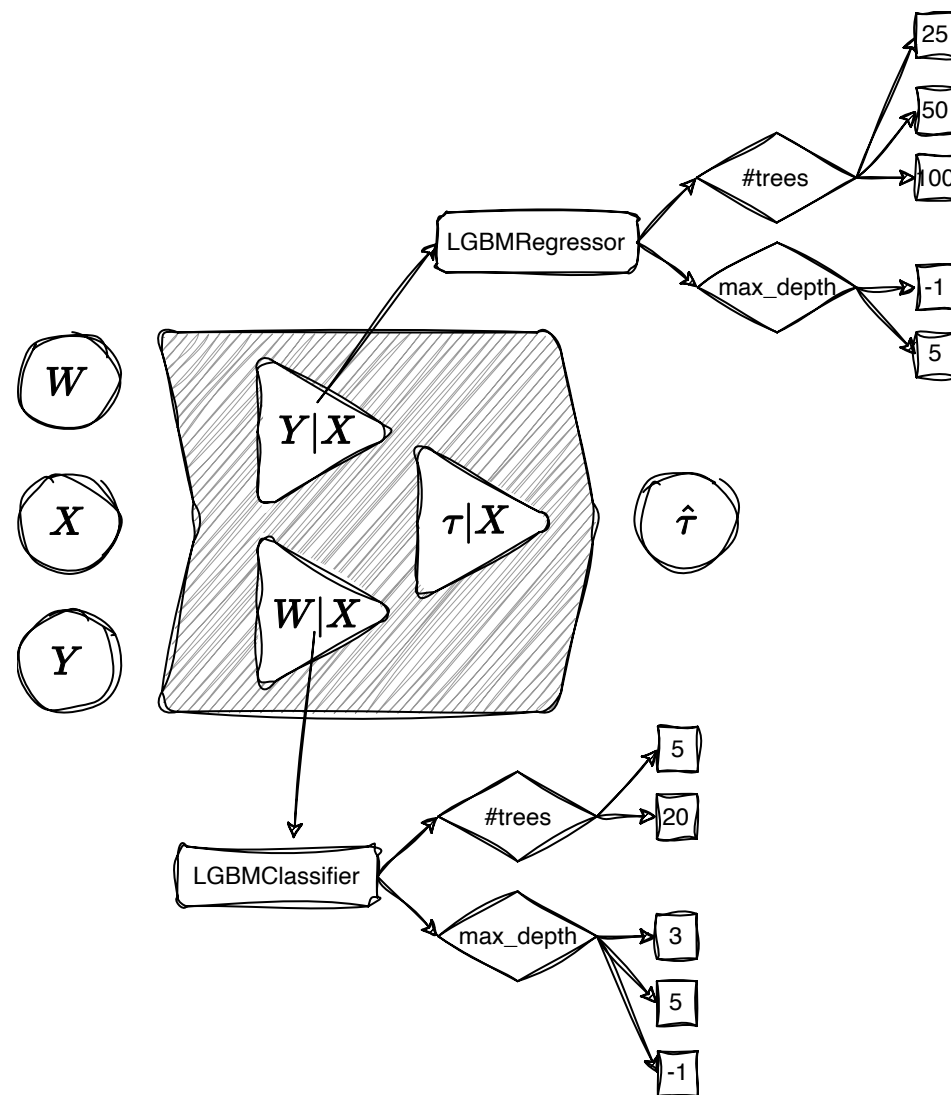
Performing a grid search



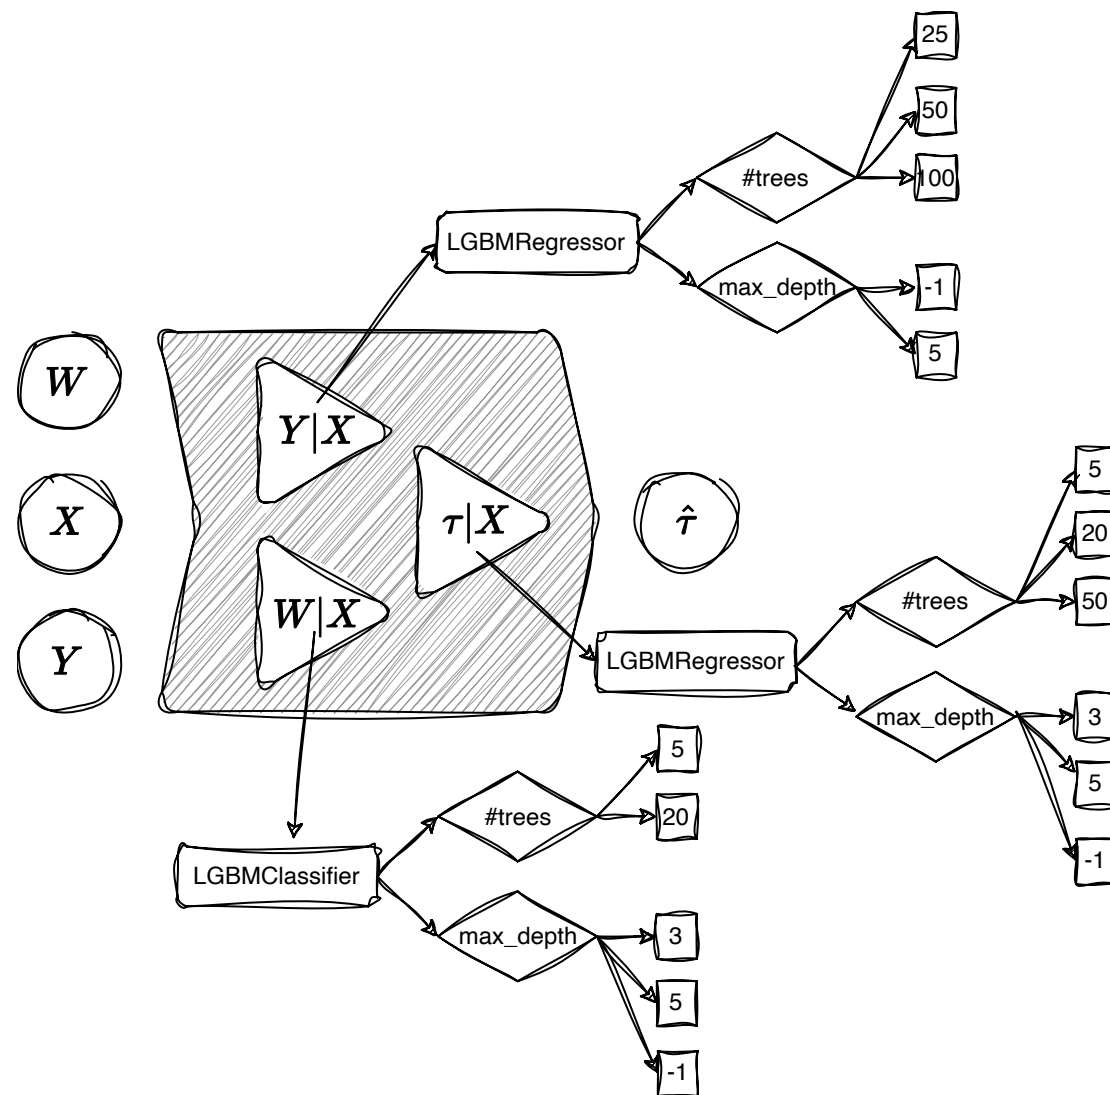
Performing a grid search



Performing a grid search



Performing a grid search



Performing a grid search

```
base_learner_grid = {  
    "outcome_model": [LGBMRegressor],  
    "propensity_model": [LGBMClassifier],  
    "treatment_model": [LGBMRegressor],  
}
```

```
param_grid = {  
    "outcome_model": {  
        "LGBMRegressor": {"n_estimators": [25, 50, 100], "max_depth": [-1, 5]}  
    },  
    "treatment_model": {  
        "LGBMRegressor": {"n_estimators": [5, 20, 50], "max_depth": [-1, 3, 5]}  
    },  
    "propensity_model": {  
        "LGBMClassifier": {"n_estimators": [5, 50], "max_depth": [-1, 3, 5]}  
    },  
}
```

Performing a grid search

```
gs = MetaLearnerGridSearch(  
    metalearner_factory=RLearner,  
    metalearner_params={"is_classification": False, "n_variants": 2},  
    base_learner_grid=base_learner_grid,  
    param_grid=param_grid,  
)
```

```
gs.fit(X_train, y_train, w_train, X_validation, y_validation, w_validation)
```

gs.results_

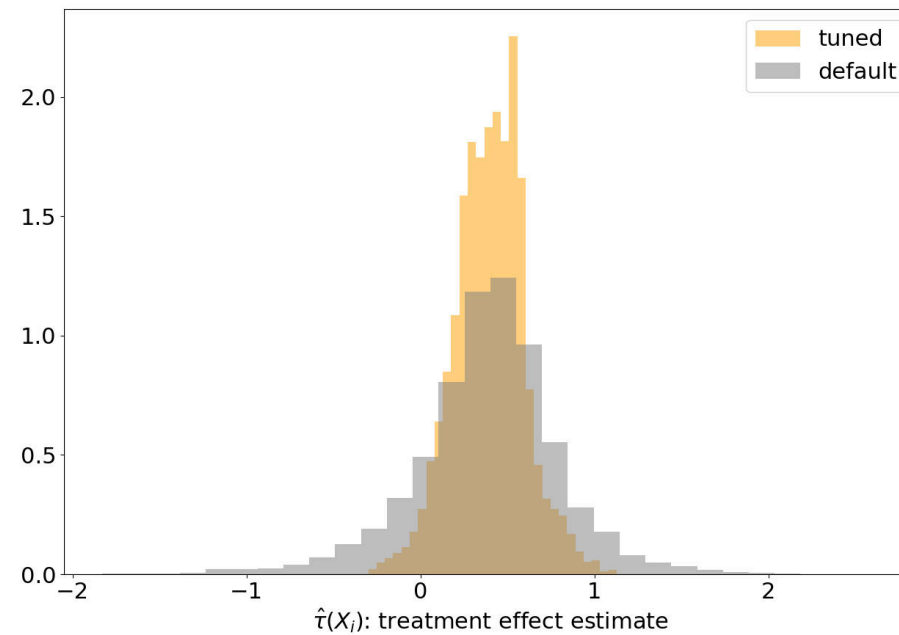
hyperparameters	time fit	time score	train propensity	train outcome	train r_loss	trea
-1, 25, -1, 5, -1, 5	0.899935	0.304743	-0.631725	-0.817461	0.795676	-1.0
-1, 25, -1, 5, -1, 20	0.965532	0.312325	-0.631725	-0.817461	0.798854	-1.0
-1, 25, -1, 5, -1, 50	1.19587	0.365287	-0.631725	-0.817461	0.804784	-1
...	
5, 25, 3, 5, 3, 20	1.79398	0.108887	-0.630564	-0.818076	0.796231	-1.0
...	
5, 100, 5, 50, 5,	1.000000	0.400000	0.000000	0.000000	0.000000	1

gs.results_

hyperparameters	time fit	time score	train propensity	train outcome	train r_loss	treat
-1, 25, -1, 5, -1, 5	0.899935	0.304743	-0.631725	-0.817461	0.795676	-1.0
-1, 25, -1, 5, -1, 20	0.965532	0.312325	-0.631725	-0.817461	0.798854	-1.0
-1, 25, -1, 5, -1, 50	1.19587	0.365287	-0.631725	-0.817461	0.804784	-1.0
...
5, 25, 3, 5, 3, 20	1.79398	0.108887	-0.630564	-0.818076	0.796231	-1.0
...
5, 100, 5, 50, 5,	1.000000	0.400000	0.000000	0.000000	0.000000	1.0

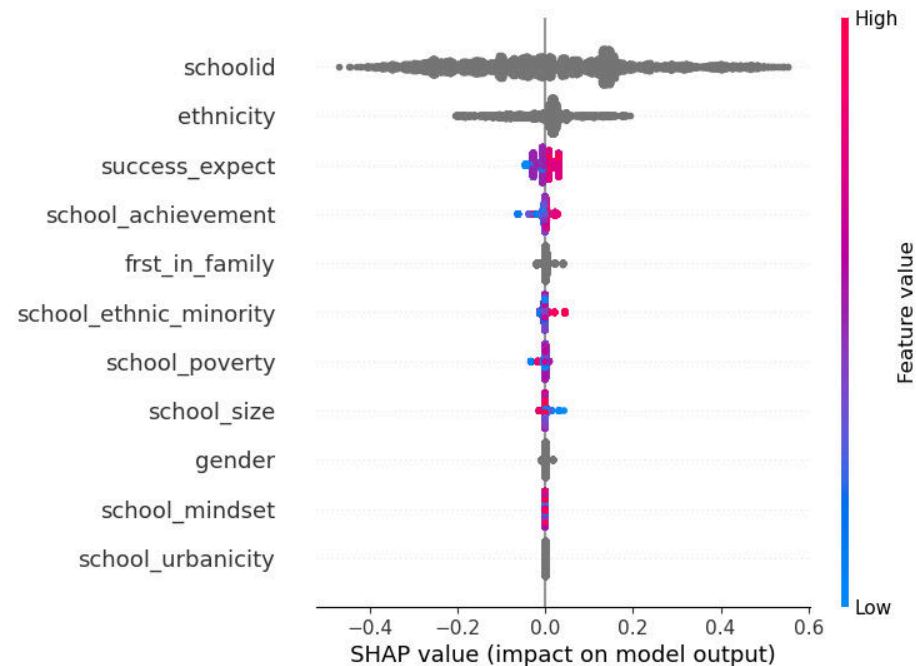
Predicting with a **tuned** MetaLearner

```
tuned_rlearner.predict(df[feature_columns], is_oos=False)
```






```
from shap import TreeExplainer, summary_plot
explainer = learner.explainer()
shap_values = explainer.shap_values(df[feature_columns], TreeExplainer)
summary_plot(shap_values[0], features=df[feature_columns])
```



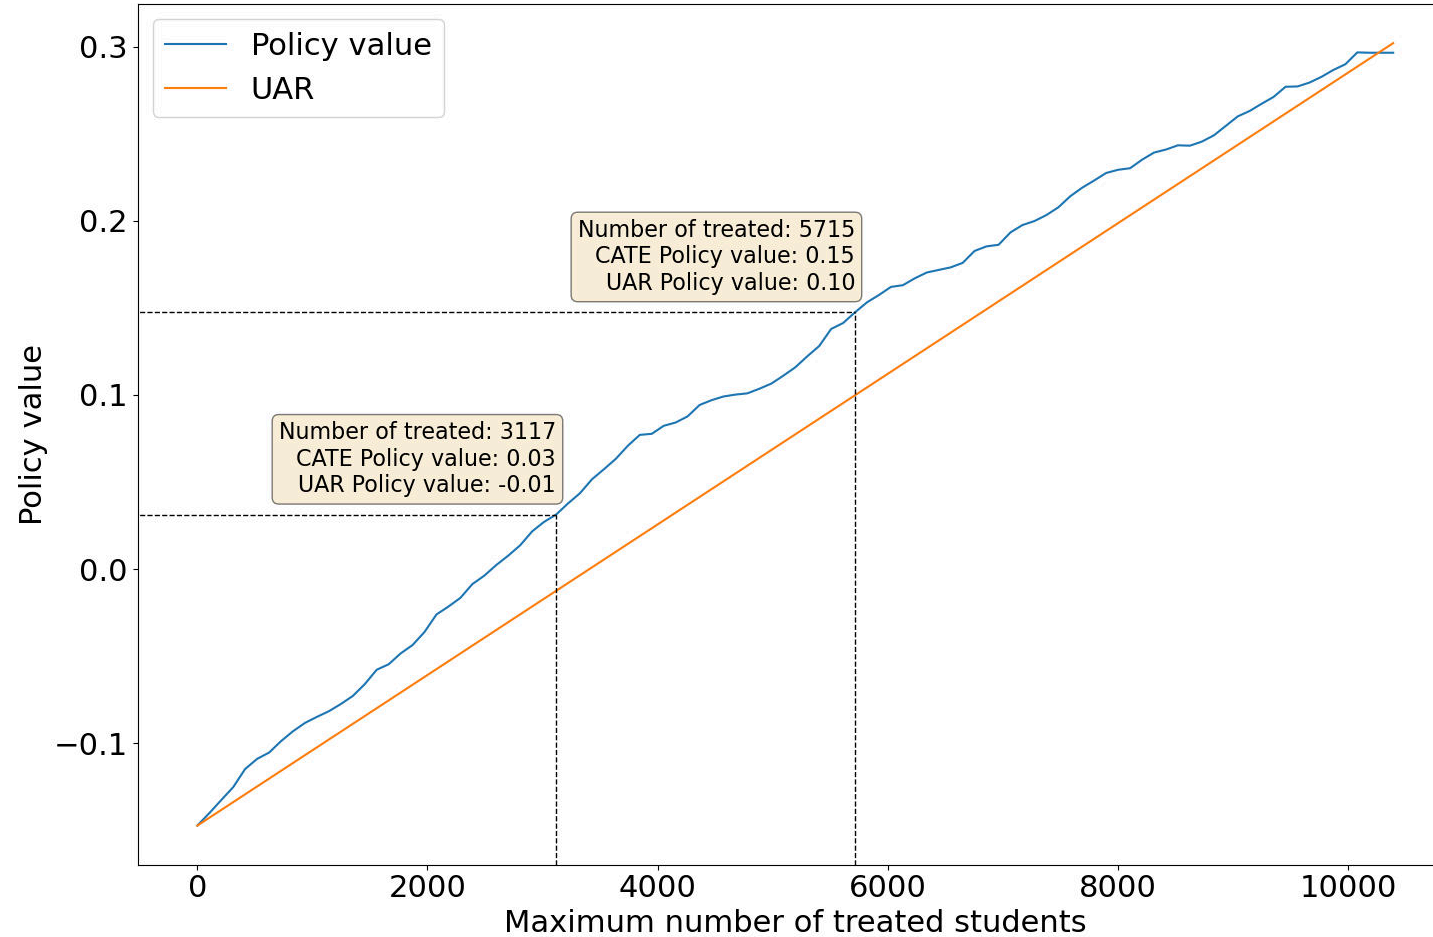
But now, how are we actually doing?

We can define the policy value as:


$$V(\pi) = \mathbb{E}[Y_i(\pi(X_i))]$$


Using our CATE estimates, we can define a policy that targets the most promising students, specifically, those with the highest CATE estimates.

But now, how are we actually doing?



Making it tangible





**Would you like to work on
such topics, too?**

Join us!

quantco.com

DEEP LEARNING

Deep Learning Engineer

HYBRID — FULL-TIME EUROPE

APPLY

Research Scientist - Virdx

HYBRID — FULL-TIME LONDON, ENGLAND / ZURICH, SWITZERLAND

APPLY

ENGINEERING

Senior Software Engineer

HYBRID — FULL-TIME EUROPE

APPLY

Software Engineer

HYBRID — FULL-TIME KARLSRUHE, BADEN-WÜRTTEMBERG

APPLY

Software Engineer

HYBRID — FULL-TIME ZURICH, SWITZERLAND

APPLY

Software Engineer

HYBRID — FULL-TIME BERLIN, BERLIN

APPLY

Software Engineer

HYBRID — FULL-TIME MUNICH, BAVARIA

APPLY

Software Engineering Intern

ON-SITE — INTERN EUROPE

APPLY

DATA SCIENCE

Data Science Intern

ON-SITE — INTERN EUROPE

APPLY

Data Scientist

HYBRID — FULL-TIME ZURICH, SWITZERLAND

APPLY

Data Scientist

HYBRID — FULL-TIME MUNICH, BAVARIA

APPLY

Data Scientist

HYBRID — FULL-TIME BERLIN, BERLIN

APPLY

Data Scientist

HYBRID — FULL-TIME LONDON, ENGLAND

APPLY

Medical Physicist - Virdx

HYBRID — FULL-TIME BOSTON, MASSACHUSETTS

APPLY

Quantitative Researcher


HYBRID — FULL-TIME EUROPE

APPLY

Quantitative Researcher

HYBRID — FULL-TIME USA

APPLY



**Please leave feedback on
GitHub! :)**

github.com/QuantCo/metalearners

github.com/kklein/sps24-metalearners

Backup

Data dictionary

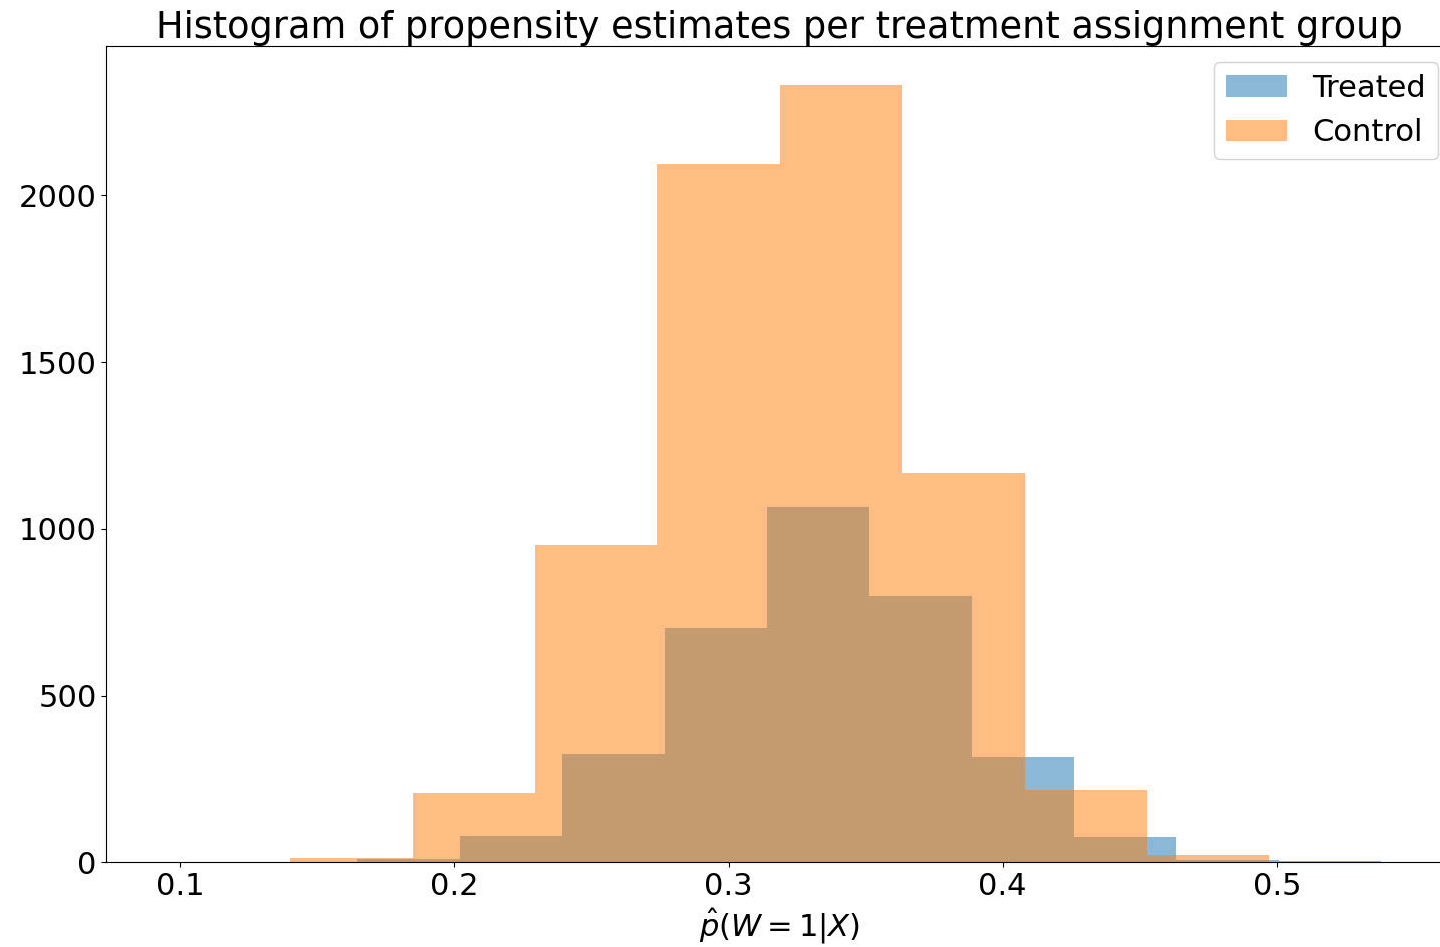
Name	Type	Meaning
ethnicity	categorical	student race/ethnicity
gender	categorical	student-identified gender
success_expect	discrete	self-reported expectations for success in the future
frst_in_family	boolean	first in family to go to college
schoolid	categorical	identifier for each of 76 high schools
school_urbanicity	categorical	school's urbanicity (urban, rural, etc.)
school_mindset	numerical	school's mean mindset

Conventional assumptions for estimating CATEs

- Positivity/overlap
- Conditional ignorability/unconfoundedness
- Stable Unit Treatment Value (SUTVA)

A randomized control trial usually gives us the first two for free.

For more information see e.g. [Athey and Imbens, 2016](#).



Policy value estimation

We estimated policy values via the 'Inverse-Propensity Weighting' estimator:

$$\hat{V}_{IPW}(\pi) = \frac{1}{n} \sum_{i=1}^n \frac{Y_i \mathbb{I}[W_i = \pi(X_i)]}{\Pr[W_i = \pi(X_i) | X_i]}$$

For more details, please see [Stefan Wager's lecture notes](#).

Python implementations of MetaLearners

	metalearners	causalml	econml
MetaLearner implementations	✓	✓	✓
Support* for pandas , scipy , polars	✓	✗	✗
HPO integration	✓	✗	✗
Concurrency across base models	✓	✗	✗
>2 treatment variants	✓	✓	✗
Classification*	✓	✗	✓
Other Causal Inference methods	✗	✓	✓