

Konrad Gawda



# Float

Everything You Wanted to Know About

# Konrad Gawda



**Python** - programmer and trainer

**Cloud** Evangelist

Videocast host



1231231231231231231



The image displays a sequence of 12 characters: 1, 2, 3, 1, 2, 3, 1, 2, 3, 1, 2, 3, 1, 2, 3, 1. A small dot is positioned directly beneath the 4th character (the first '1' after the initial '123'). Below the entire sequence, a horizontal double-headed arrow spans from the vertical position of the 4th character to the vertical position of the 10th character.

```
import numbers
```

```
numbers.Number
```



```
numbers.Complex → numbers.Real → numbers.Rational → numbers.Integral
```



```
float
```

1.0


1e0

float('1.0')

float(1)

1/1

# Binary representation



## THE ROYAL SHIP TO CANADA

On September 25th, Their Royal Highnesses Princess Elizabeth and the Duke of Edinburgh are due to leave Liverpool in the Canadian Pacific liner "Empress of France" for Quebec, whence they will start their tour of the Dominion. A special suite on "A" deck has been reserved for them and their entourage. This fine liner, depicted here, has a gross weight of 26,400 tons, and is 601 ft. long. Her geared turbine engines, driving two propellers, give her a speed of 18 knots. This drawing was made with the kind co-operation and assistance of the Canadian Pacific Railways.

### KEY TO FEATURES

- (1) The Royal Suite. (2) Duke of Edinburgh's staff cabins. (3) Princess Elizabeth's staff cabins. (4) Princess Elizabeth's Standard from mainmast. (5) Radar scanner and cabin. (6) Wheelhouse and navigating bridge. (7) Radio lead-in. (8) Lift machinery. (9) Funnel uptakes from aft boiler room. (10) Engine room hatch and ventilators. (11) Chimney to the Engine room on the promenade deck. (12) First class lounge. (13) Lift and stairs to all decks. (14) Long Gallery leading to Cocktail Bar under bridge. (15) Enclosed Promenade deck. (16) Officers' cabins. (17) Bath rooms. (18) First class cabins. (19) Crew and stewards' quarters. (20) Cargo holds. (21) Forward boiler room (4 boilers). (22) Oil fuel tanks. (23) Aft boiler room (4 boilers). (24) First class dining room. (25) Captain's table, at which Princess Elizabeth and the Duke of Edinburgh may occasionally dine. (26) Main kitchen serving both dining rooms. (27) Geared turbine engine room. (28) Diesel-electricity generator. (29) Port propeller shaft tunnel. (30) Cold store. (31) Tourist class entrance hall. (32) Tourist lounge. (33) Tourist cabins. (34) Tourist dining room. (35) Tourist cabins. (36) Port and starboard propellers. (37) Stern anchor capstan. (38) Docking bridge.

*Almost all platforms map Python floats to IEEE 754 binary64  
“double precision” values*

<https://docs.python.org/3/tutorial/floatingpoint.html>

```
typedef struct {  
    PyObject_HEAD  
    double ob_fval;  
} PyFloatObject;
```

<https://github.com/python/cpython/blob/main/Include/cpython/floatobject.h>

```
import sys, ctypes
f = 1.0

sys.getsizeof(f)
24

ctypes.string_at(id(f), sys.getsizeof(f)).hex()
'03000000000000000059950000000000000000000000f03f'

ctypes.string_at(id(f), sys.getsizeof(f))[-8:].hex()
'000000000000f03f'

bin(int.from_bytes(ctypes.string_at(id(f), sys.getsizeof(f))[-8:]))
'0b1111000000111111'

bin(int.from_bytes(ctypes.string_at(id(f), sys.getsizeof(f))[-8:],\
byteorder=sys.byteorder))
'0b1111111111000000000000000000000000000000000000000000000000000000'

bin(int.from_bytes(ctypes.string_at(id(f), sys.getsizeof(f))[-8:],\
byteorder=sys.byteorder)).replace('0b', '').zfill(64)
'0011111111110000000000000000000000000000000000000000000000000000'
```



$0.625_{10}$

$0.a_{16}$

$0.101_2$

$$9\text{e}13 = 9 \times 10^{**} 13$$

$$1.625\text{e}6 = 1.625 \times 10^{**} 6$$

*Rather not: 1625e3, 0.1625e7*

$$1.625\text{e}-6 = 1.625 \times 10^{**} -6 = 1.625 / 10^{**} -6 \quad x^{-n} = 1 / x^n$$

*Time for binary numbers*

$$1.101_p 101 = 1.101_2 \times 2^{**} 101_2$$


$$1.625 * 2^{**} 5 = 52$$

$$1.101_p -101 = 1.101_2 \times 2^{**} -101_2$$

$$1.101 \ 11 \rightarrow 1.101_2 \times 2^{**} (11-1000)_2$$

$$(4).101 \ 11 \rightarrow 1.101_2 \times 2^{**} (11-1000)_2$$

# IEEE 754 binary64 ("double")



$$(-1)^{sign} \times 2^{exp - 1023} \times 1.\textit{fraction}$$

- Sign bit: 1 bit
- Exponent: 11 bits
- Significand / fraction: 53 bits (52 explicitly stored)

# IEEE 754...



**William "Velvel" Kahan**

A primary architect of the Intel 80x87 floating-point coprocessor and IEEE 754 floating-point standard.

# Exponent

$$(-1)^{sign} \times 2^{\mathbf{exp} - 1023} \times 1.\mathit{fraction}$$

- $000_{16}$ 
  - zero (when fraction == 0)
  - *subnormal (denormalized) numbers*
- $001_{16} - 7fe_{16}$ 
  - $001 \rightarrow 2^{*-2022}_{10}$
  - $3ff (1023_{10}) \rightarrow 2^{*0}$
  - $7fe (2046_{10}) \rightarrow 2^{*2023}_{10}$
- $7ff_{16}$ 
  - infinity (when fraction == 0)
  - NaN

000000000000

000000000001


011111111111

111111111110


111111111111



 1.0 0.5 0.125 0.1

 -0.1

# Programmer's point of view





# float methods

```
float.is_integer()
```

```
float.as_integer_ratio()
```

```
(0.5).as_integer_ratio() → (1, 2)
```

```
(0.1).as_integer_ratio() → (3602879701896397, 36028797018963968)
```

*So why Python does not print 1/10 as 0.1000000000000000005551115123 ?*

Float is by default printed rounded to *machine precision*.

# float methods

`float.hex()`

`float.fromhex(s)` # *classmethod*

e.g. `'-0x1.8000000000000p+0'`

[sign] ['0x'] integer ['.' fraction] ['p' exponent]

*exponent:  $2^{**exp_{10}}$*

IEEE 754

C's %a format  
character

Java's  
`Double.toHexString`

'0x0.0p+0'

0.0

'0x1.00000000000000p+0'

1.0

'0x1.c0000000000000p+2'

7.0

'0x1.f8000000000000p+5'

63.0

'0x1.ff800000000000p+9'

1023.0

'0x1.00000000000000p+10'

1024.0

$$1c_{16} = 28_{10}$$

'0x1.00000000000000p-1'

0.5

'0x1.00000000000000p-3'

0.125

'0x1.99999999999999ap-4'

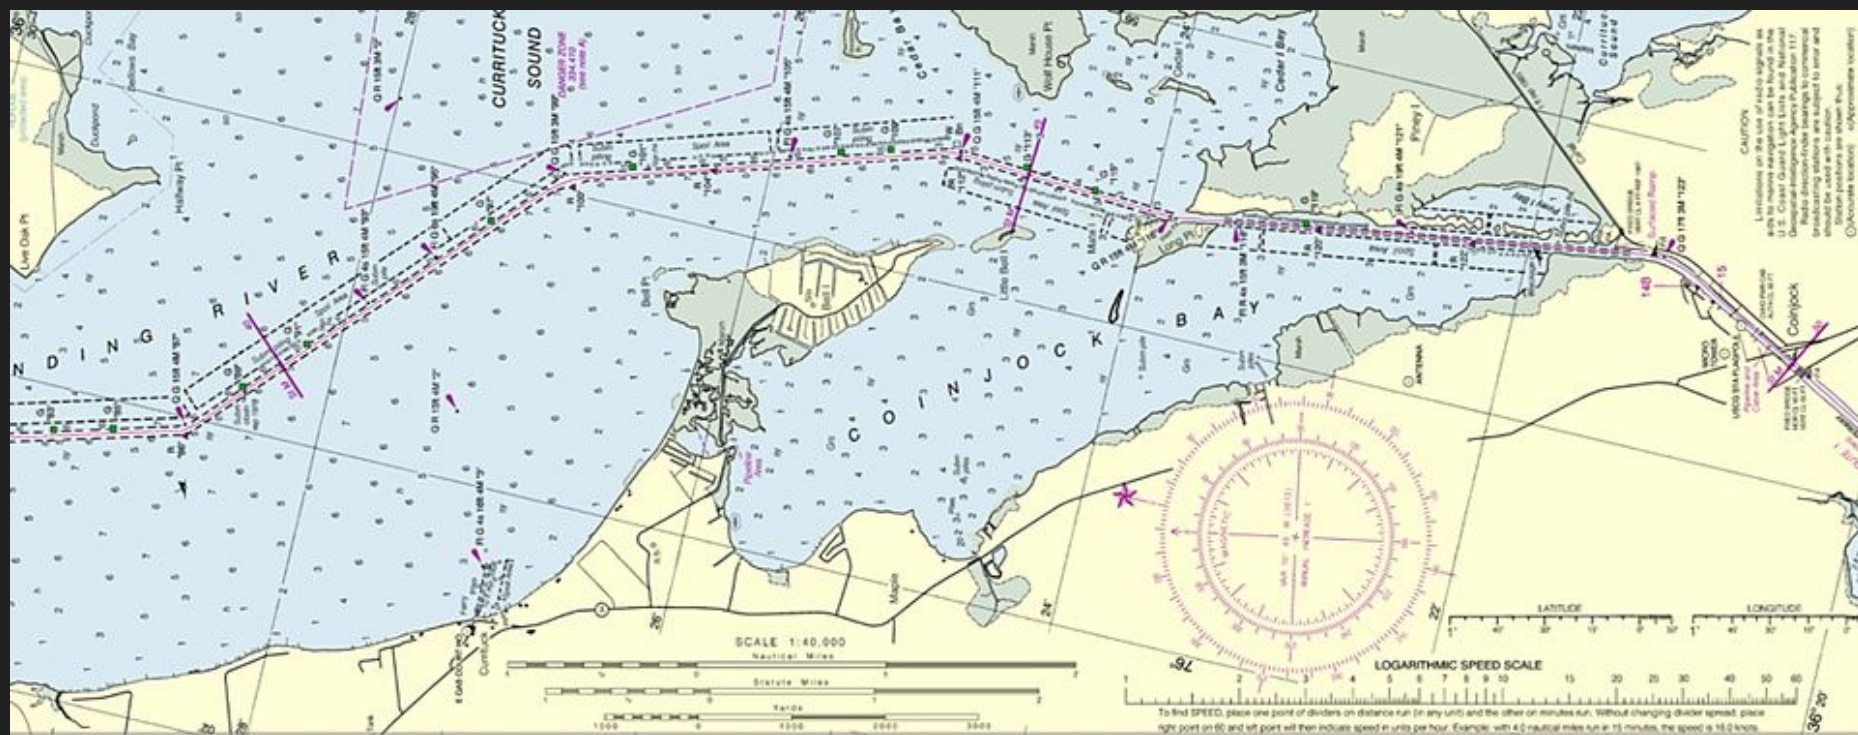
0.1

'-0x1.99999999999999ap-4'

-0.1

$$9_{16} = 1001_2$$

# Precision



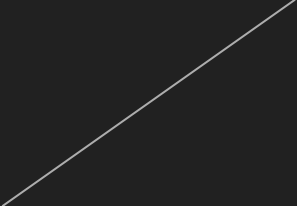
`sys.float_info`

```
sys.float_info(  
    max=1.7976931348623157e+308,  
    max_exp=1024,  
    max_10_exp=308,  
    min=2.2250738585072014e-308,  
    min_exp=-1021,  
    min_10_exp=-307,  
    dig=15,  
    mant_dig=53,  
    epsilon=2.220446049250313e-16,  
    radix=2,  
    rounds=1  
)
```

1231231231231231231



September 2022  
CVE-2020-10735: Prevent  
DoS by large int<->str  
conversions



```
sys.float_info.dig → 15  
sys.get_int_max_str_digits() → 4300
```

`math.ulp(x)` - *Unit in the Last Place* (since Python 3.9)



```
for x in range(100):  
    print(f"ulp(2**{x}) = {math.ulp(2**x)}")
```

```
ulp(2**0) = 2.220446049250313e-16  
ulp(2**1) = 4.440892098500626e-16  
ulp(2**2) = 8.881784197001252e-16  
ulp(2**3) = 1.7763568394002505e-15  
ulp(2**4) = 3.552713678800501e-15  
ulp(2**5) = 7.105427357601002e-15  
ulp(2**6) = 1.4210854715202004e-14  
ulp(2**7) = 2.842170943040401e-14  
ulp(2**8) = 5.684341886080802e-14  
ulp(2**9) = 1.1368683772161603e-13  
ulp(2**10) = 2.2737367544323206e-13  
ulp(2**11) = 4.547473508864641e-13  
ulp(2**12) = 9.094947017729282e-13  
ulp(2**13) = 1.8189894035458565e-12  
ulp(2**14) = 3.637978807091713e-12  
ulp(2**15) = 7.275957614183426e-12  
ulp(2**16) = 1.4551915228366852e-11  
ulp(2**17) = 2.9103830456733704e-11  
ulp(2**18) = 5.820766091346741e-11  
ulp(2**19) = 1.1641532182693481e-10  
ulp(2**20) = 2.3283064365386963e-10
```

```
ulp(2**21) = 4.656612873077393e-10  
ulp(2**22) = 9.313225746154785e-10  
ulp(2**23) = 1.862645149230957e-09  
ulp(2**24) = 3.725290298461914e-09  
ulp(2**25) = 7.450580596923828e-09  
ulp(2**26) = 1.4901161193847656e-08  
ulp(2**27) = 2.9802322387695312e-08  
ulp(2**28) = 5.960464477539063e-08  
ulp(2**29) = 1.1920928955078125e-07  
ulp(2**30) = 2.384185791015625e-07  
ulp(2**31) = 4.76837158203125e-07  
ulp(2**32) = 9.5367431640625e-07  
ulp(2**33) = 1.9073486328125e-06  
ulp(2**34) = 3.814697265625e-06  
ulp(2**35) = 7.62939453125e-06  
ulp(2**36) = 1.52587890625e-05  
ulp(2**37) = 3.0517578125e-05  
ulp(2**38) = 6.103515625e-05  
ulp(2**39) = 0.0001220703125  
ulp(2**40) = 0.000244140625  
ulp(2**41) = 0.00048828125
```

```
ulp(2**42) = 0.0009765625  
ulp(2**43) = 0.001953125  
ulp(2**44) = 0.00390625  
ulp(2**45) = 0.0078125  
ulp(2**46) = 0.015625  
ulp(2**47) = 0.03125  
ulp(2**48) = 0.0625  
ulp(2**49) = 0.125  
ulp(2**50) = 0.25  
ulp(2**51) = 0.5  
ulp(2**52) = 1.0  
ulp(2**53) = 2.0  
ulp(2**54) = 4.0  
ulp(2**55) = 8.0  
ulp(2**56) = 16.0  
ulp(2**57) = 32.0  
ulp(2**58) = 64.0  
ulp(2**59) = 128.0  
ulp(2**60) = 256.0  
ulp(2**61) = 512.0  
...
```

*Any float  $x$  within  $2^{52} - 2^{53}$  has int precision*









































4 503 599 627 370 496 ... 9 007 199 254 740 992

4.503600e+15 ... 9.007199e+15

0x1.0000000000000p+52 ... 0x1.0000000000000p+53

4.5 ... 9 × {Ger./Fr./It./Pol. **billiard\***, Eng. **quadrillion**, SI: **Peta**}

# Special symbols

					
<b>A-Alfa</b> Diver Down Keep Clear	<b>B-Bravo</b> Dangerous Cargo	<b>C-Charlie</b> Yes	<b>D-Delta</b> Keep Clear	<b>E-Echo</b> Altering Course to Starboard	<b>F-Foxtrot</b> Disabled
					
<b>G-Golf</b> Want a Pilot	<b>H-Hotel</b> Pilot on Board	<b>I-India</b> Altering Course to Port	<b>J-Juliet</b> On Fire Keep Clear	<b>K-Kilo</b> Desire to Communicate	<b>L-Lima</b> Stop Instantly
					
<b>M-Mike</b> I Am Stopped	<b>N-November</b> No	<b>O-Oscar</b> Man Overboard	<b>P-Papa</b> About To Sail	<b>Q-Quebec</b> Request Practique	<b>R-Romeo</b> Do Not Pass Ahead of Me
					
<b>S-Sierra</b> Engines Going Astern	<b>T-Tango</b> Keep Clear	<b>U-Uniform</b> Standing Into Danger	<b>V-Victor</b> Require Assistance	<b>W-Whiskey</b> Require Medical Assistance	<b>X-Xray</b> Stop Your Intention
					
<b>Y-Yankee</b> Am Dragging Anchor	<b>Z-Zulu</b> Require A Tug	0	1	2	3
					
4	5	6	7	8	9
					
<b>1st Repeater</b>	<b>2nd Repeater</b>	<b>3rd Repeater</b>	<b>Code &amp; Answering</b>		

# Exponent

- $000_{16}$ 
  - zero (when fraction == 0)
  - *subnormal (denormalized) numbers*
- $001_{16} - 7fe_{16}$ 
  - $001 \rightarrow 2^{*-2022}_{10}$
  - $3ff (1023_{10}) \rightarrow 2^{*0}$
  - $7fe (2046_{10}) \rightarrow 2^{*2023}_{10}$
- $7ff_{16}$ 
  - infinity (when fraction == 0)
  - NaN

000000000000

000000000001

011111111111

111111111110

111111111111

# This is still *float*

IEEE 754


```
float('Inf')    # inf, Infinity, ...  
float('-Inf')   # -inf, -Infinity, ...  
math.inf, -math.inf  
math.isinf(x)
```

```
-0.0, 0/-1, -1/float('inf')  
0.0 == -0.0  
-0.0 + 0.0 == 0.0
```

```
float('NaN')    # nan, -nan, ... - Not a Number  
math.nan  
math.isnan(x)  
bool(math.nan) == True
```

Comparisons like `=`, `<`, `>=`, ... with NaN - all return False

# Processor modes



# FLT\_ROUNDS (float.h macro)

`sys.float_info.rounds`

integer representing the rounding mode for floating-point arithmetic. This reflects the value of the system FLT\_ROUNDS macro at interpreter startup time:

- -1: indeterminable
- 0: toward zero
- 1: to nearest ("unbiased") - default
- 2: toward positive infinity
- 3: toward negative infinity

IEEE 754



```
from ctypes.util import find_library

libm = ctypes.cdll.LoadLibrary(find_library('m'))

FE_TONEAREST = 0x0000
FE_DOWNWARD = 0x0400
FE_UPWARD = 0x0800
FE_TOWARDZERO = 0x0c00

libm.fesetround(FE_TONEAREST)
(1/10).hex()      → '0x1.999999999999ap-4'

libm.fesetround(FE_TOWARDZERO)
(1/10).hex()      → '0x1.9999999999999p-4'
```



```
libm.fesetround(FE_TONEAREST)
```

9007199254740992.0 + 1

9007199254740992.0

```
libm.fesetround(FE_UPWARD)
```

9007199254740992.0 + 1

9007199254740994.0

# round() - processor-independent

round() - x rounded to n digits, rounding halves **to even**.

round(0.5)

0

round(1.5)

2

round(2.5)

2

round(3.5)

4

a // b

Int division, rounded towards *-Infinity*.

Result can be float

int(), math.trunc()

truncates (towards zero)

math.floor(), x.\_\_floor\_\_

math.ceil(), x.\_\_ceil\_\_

round  
f"{a:.1f}"

a // 1  
math.floor

int  
math.trunc

math.ceil



# Traps



# Trap

```
f = 0
while f != 1:
    f += 0.1
```

Do you really need "equal"?

```
f = 0
while f < 1:
    f += 0.1
```

# Trap

```
0.1 + 0.1 + 0.1 == 0.3
```

```
False
```

```
math.fsum([0.1, 0.1, 0.1]) == 0.3
```

```
True
```

# Trap - ?

```
#include <stdio.h>
int main() {
    float meters = 0;
    int iterations = 100000000;
    for (int i = 0; i < iterations; i++) {
        meters += 0.01;
    }
    printf("Expected: %f km\n", 0.01 * iterations / 1000 );
    printf("Got: %f km \n", meters / 1000);
}
```

```
Expected: 10000.000000 km
Got: 262.144012 km
```

# Trap

`13.716 / 4.572 → 3.0`

`13.716 % 4.572 → 4.571999999999999`

`13.716 // 4.572 → 2.0`

or...

`divmod(13.716, 4.572)`  
`(2.0, 4.571999999999999)`

*divmod(a, b) - For integers, the result is the same as (a // b, a % b). For floating point numbers the result is (q, a % b), where q is usually math.floor(a / b) but may be 1 less than that.*



# Real world examples



# Patriot, 1991

Dhahran, Saudi Arabia. Patriot did not shoot at Scud rocket.

Internal clock gives ticks each 0.1 s.

Increments by 0.1 in 24 bit *fixed point* register. This gives error of 0.000000095.

With 100h of standby this gives error of 0.34 s.



Other Scud rocket, that was shot down

# Ariane 5, 1996

First Ariane 5 flight.

Cast error (64-bit float -> 16-bit signed int). Ada language.

```
L_M_BV_32 := TBD.T_ENTIER_32S ((1.0/C_M_LSB_BV) * G_M_INFO_DERIVE(T_ALG.E_BV));
```

```
if L_M_BV_32 > 32767 then
    P_M_DERIVE(T_ALG.E_BV) := 16#7FFF#;
elsif L_M_BV_32 < -32768 then
    P_M_DERIVE(T_ALG.E_BV) := 16#8000#;
else
    P_M_DERIVE(T_ALG.E_BV) := UC_16S_EN_16NS(TDB.T_ENTIER_16S(L_M_BV_32));
end if;
```

```
P_M_DERIVE(T_ALG.E_BH) :=
    UC_16S_EN_16NS (
        TDB.T_ENTIER_16S ((1.0/C_M_LSB_BH) *
            G_M_INFO_DERIVE(T_ALG.E_BH))
    );
```

# Sleipner A, 1991

The concrete foundations for an oil rig collapsed during the test.

The incorrect calculations attributed to the NASTRAN program (in Fortran) - no details.

<https://github.com/nasa/NASTRAN-95>



# Schleswig-Holstein parliamentary elections, 1992

The electoral threshold was 5%.

The Green Party received 4.97%.

The result was rounded to 5.0% on printout.

The program has been used for years.



Reisemagazin Online


# Kerbal Space Program: Deep Space Kraken

## Minecraft: Far Lands





# Similar species



# Decimal

```
from decimal import Decimal
```

```
Decimal('7.325').quantize(Decimal('0.01'), rounding=ROUND_DOWN)
```

```
json.loads('0.1', parse_float=Decimal)
```



# from fractions import Fraction

```
Fraction.limit_denominator(max_denominator=1000000)
```

```
Fraction(1.1).limit_denominator()
```

```
Fraction(11, 10)
```

```
class Fraction(numbers.Rational):  
    def __new__(cls, numerator=0, denominator=None, *, _normalize=True):  
        self = super(Fraction, cls).__new__(cls)  
        ...  
        self._numerator = numerator  
        self._denominator = denominator  
        return self
```

<https://github.com/python/cpython/blob/main/Lib/fractions.py>

# mpmath

```
>>> from mpmath import mp
>>> mp.dps = 50
>>> print(mp.quad(lambda x: mp.exp(-x**2), [-mp.inf, mp.inf]) ** 2)
3.1415926535897932384626433832795028841971693993751
```

# PySim

```
>>> from sympy import *  
>>> x, t, z, nu = symbols('x t')  
>>> integrate(sin(x**2), (x, -oo, oo))
```

$$\frac{\sqrt{2} \cdot \sqrt{\pi}}{2}$$

uses mpmath :)

# Thanks for your attention



Konrad Gawda

[linkedin: konradgawda](#)

